

How Much Information Fits in a Vector?

Anonymous authors

Paper under double-blind review

Abstract

Recent work in neural network interpretability has suggested that hidden activations of some deep models can be viewed as linear projections of much higher-dimensional vectors of sparse latent “features.” In general, this kind of representation is known as a superposition code. This work presents an information-theoretic account of superposition codes in a setting applicable to interpretability. We show that, when the number k of active features is very small compared to the number N of total possible features, surprisingly simple methods can decode these representations provided that d is a constant factor greater than the Shannon limit. Specifically, when $\ln k / \ln N \leq \eta < 1$ and H is the entropy of the latent vector in bits, it suffices that $d/H \geq C(\eta)$ for a certain increasing function $C(\eta)$. However, the optimal constant $C(\eta)$ varies significantly depending on the decoding method used. For example, we show that the popular “top- k ” method typically requires a factor of $C = 4$, whereas we introduce a family of cheap iterative methods (matching- k) that can achieve around $C = 2$. Our results address a lack of practical references on superposition codes in a very sparse regime.

1 Introduction

If each neuron in a given neural network coded for a “meaningful” feature of its input, we could hope to reverse-engineer this network’s overall behavior on a neuron-by-neuron basis. However, individual neurons of real-world networks often lack clear interpretations. For example, both language models and vision models have been found to learn neurons that correlate simultaneously with apparently unrelated features. (See for example Nguyen et al. (2016), Zhang & Wang (2023) and Olah et al. (2020).)

The difficulty of interpreting a network in terms of its local activity—and in particular, the appearance of so-called “polysemantic neurons”—is not surprising from a connectionist viewpoint. Since at least the 1980s, proponents of neural networks have argued that these systems may naturally use **distributed representations**—coding schemes where individual features are represented by patterns spread over many neurons, and conversely where each neuron carries information on many features. (This term was apparently coined in Rumelhart et al. (1986), Chapter 3.) In contrast, a *local* representation would dedicate each neuron to a




Figure 1: A coarse code representing a point on a plane. Each “neuron,” drawn as a red or blue square, encodes whether the point belongs to an associated “receptive field.” Although no neuron gives specific information on the position of the point, the overall code determines its position with reasonable accuracy.

single feature. See Thorpe (1989) for a general discussion of local and distributed codes. Figure 1 illustrates a classic example of a coarse code, one kind of distributed representation.

As of now, relatively little is known about how deep neural networks learn to represent information in their hidden layers or to what extent this information can be interpreted. However, should “interpretable features” exist, the connectivist viewpoint makes it natural that they would typically be stored with non-local codes. This is a common assumption in interpretability research today; for example, when Meng et al. (2022) intervened on an MLP layer of a language model to “edit” a factual association, both the “subject” and the “fact” were modeled as vectors of neuron activations rather than as individual neurons.

How can we infer latent features learned by a neural network? One simple proposal is to model an activation vector x as a linear projection

$$x = Fy$$

of some high-dimensional and *sparse* vector y of latent features. We refer to the columns of F as codewords and the whole matrix F as a dictionary. Since x is a linear superposition of codewords, we will call it a **superposition code** for y . It is a remarkable fact that, for certain matrices F , the linear projection $x = Fy$ can code uniquely for a sparse vector y even when the dimension of x is (logarithmically) smaller than the dimension of y . The task of inferring the sparse vector from its linear projection is known as sparse reconstruction, and the task of inferring the dictionary F from a distribution over x is called dictionary learning. Both of these problems have been studied in the field of compressive sensing, although with different applications in mind. (See Elad (2010) for a review of classic work in the context of signal and image processing.)

Already in 2015, Faruqui et al. (2015) used a dictionary learning method to derive sparse latent codes for word embeddings and argued that these latents were more interpretable than the original embedding dimensions. More recently, a series of works beginning with Yun et al. (2021) have applied dictionary learning to the internal representations of transformer-based language models. Cunningham et al. (2023) suggested the use of **sparse autoencoders** (SAEs) and Templeton et al. (2024); Gao et al. (2024) scaled sparse autoencoders to production-size large language models.

To infer a latent representation $y \in \mathbb{R}^N$ from an activation vector $x \in \mathbb{R}^d$, sparse autoencoders use a simple estimate of the form $\hat{y}(x) = \sigma(Gx)$ where $G: \mathbb{R}^{N \times d}$ is some learnable matrix and σ is a fixed “thresholding” operation. In Gao et al. (2024), which is representative of large-scale sparse autoencoder applications today, N was around one million and the number of non-zero coefficients k in each latent vector $\hat{y}(x)$ was around one hundred.

Templeton et al. (2024) showed that latent features learned by SAEs are often highly interpretable, and that intervention on these features allows “steering” language models in predictable ways. However, as reported in Gao et al. (2024), even SAEs with extremely large numbers of latents suffer from an apparently irreducible reconstruction error. According to Sharkey et al. (2025), understanding the limitations of SAEs—and dictionary learning in general—is an important open question in the research program of mechanistic interpretability.

Contributions

Over the course of computation, it is common for programs to use more memory than was required to store their input data. Similarly, it is natural to expect that the entropy of a “simple” (or “interpretable”) description for the residual vectors within a language model is not bounded by the entropy of language. Methods like sparse autoencoders aim to provide descriptions of these hidden activations, but relatively little is known about the nature of the information is being carried. One basic question is *how much* information can hypothetically be stored in an activation vector using a superposition code.

In many classical situations, this question is answered by well-known ideas from information theory. Given a “channel” with certain characteristics—for example, a band-limited telephone connection or a binary storage device with a certain failure rate—the amount of information we can encode is asymptotically characterized by a “channel capacity” measured in bits per unit of channel usage. The goal of this paper is to study the




Figure 2: An overview of the minimum codeword dimension d required for two different methods to reliably decode a uniformly chosen k -sparse subset of $\{1, \dots, 2^{20}\}$ from a superposition of codewords from an i.i.d. Rademacher dictionary. The inverse “bitrate” $d/H(k)$, where $H(k) = \log_2 \binom{N}{k} \approx k \log_2(eN/k)$, is indicated by the right axis. Top- k is a method currently used by sparse autoencoders, and matching pursuit is a relatively cheap iterative algorithm.

carrying capacity of sparse superposition codes in a regime that is meaningful to sparse autoencoders. Our analysis focuses on a toy example, described in more detail in Section 2.

Our main findings are the following.

1. In the asymptotic regime where $\ln k \leq \eta \ln N$ for $\eta < 1$, we prove that very simple “one-step methods” used by sparse autoencoders can reliably decode superposition codes when d is a constant factor $C(\eta)$ greater than the Shannon limit. Moreover, we give concrete lower bounds on the effective capacity $C(\eta)$ which agree well with empirical simulation.
2. On the other hand, we show that $C(\eta)$ is surprisingly large for values of (N, k) typical in applications of sparse autoencoders. In particular, we show empirically that there is a large gap between the

Second, we show that relatively cheap iterative methods

How “efficient,” in terms of bitrate, are the codes used by real neural networks? Of course, it would not make sense for a network to use a code that requires a costly iterative decoding process before it can be used. However, the success of matching pursuit suggests that neural networks may be able to improve over the bitrates of one-step estimates while paying a relatively small computational price. Although our analysis is restricted to a toy scenario, we hope these results inform future work on modeling distributed representations.

Related Work

Recently, various authors have studied the underperformance of SAEs and proposed ways to improve these methods. For example, Rajamanoharan et al. (2024) proposed *gated SAEs* to mitigate “feature shrinkage,” and Bussmann et al. (2025) proposed *Matryoska SAEs* to deal with problems related to “feature absorption.”

Especially relevant to this work is the proposal of *inference-time optimization* (ITO), which involves replacing the encoder of an SAE with an iterative optimization method at inference time—that is, after the dictionary F has already been learned. For example, Engels et al. (2024) evaluated gradient pursuit as an inference-time optimization method. However, if the restrictive encoders used by SAEs at training time cannot discern read certain attributes of the sparse latent signal, then the corresponding codewords will not appear in the learned dictionary. This may explain the limited improvements attained so far by ITO. We hope the present work helps clarify these questions, especially for readers who may not be familiar with ideas from compressive sensing.

2 Superposition Codes and Sparse Recovery

We begin by describing the toy scenario to be studied. Given a large number N , consider a map F that represents each subset $x \subseteq [N] = \{1, \dots, N\}$ by a linear combination

$$y(x) = \sum_{i \in x} f_i \in \mathbb{R}^d,$$

where the vectors $\{f_i \in \mathbb{R}^d : i \in [N]\}$ are chosen in advance and where the dimension d of the encoding is expected to be much smaller than N . The vectors f_i are called codewords for the elements of $[N]$, the collection $\{f_1, \dots, f_N\}$ is called a dictionary, and the image Fx is called a superposition code. It is useful to view x as a vector in $\{0, 1\}^N$ with coefficients

$$x_i = \begin{cases} 1 & : i \in x \\ 0 & : \text{otherwise} \end{cases}$$

and view F the matrix of column vectors $[f_1 \dots f_N]$, so that $y = Fx$ is a linear equation. In this work, we'll model our subset as a random variable X uniformly distributed over all subsets of some fixed size $k \ll N$. Keeping with the orders of magnitude discussed in Gao et al. (2024), we consider $N \approx 1000,000$ and $k \approx 100$.

The problem of recovering a sparse signal x from a linear projection Fx is known as sparse recovery. This problem is considered in the subject of compressive sensing. In this work, we will use language from coding theory and think of x as some data to be encoded and the projection $y = Fx$ as a code. Thus, sparse recovery becomes the task of “decoding” x . Note that this is opposite to the convention used in sparse autoencoders, where the map estimating x from y is called the “encoder.” (In classical applications of autoencoders, the learned representation is understood as an efficient “code” for the data distribution.)

Due to a connection with sparse linear regression, the subproblem of identifying the support of x from the projection Fx is often called model selection. When the coordinates of x take values in $\{0, 1\}$, sparse recovery reduces to a problem of model selection. In general, note that model selection presents the main difficulty of sparse recovery; once the support of x identified, its non-zero coordinates can be estimated easily e.g. by solving a least squares problem.

It remains to choose the dictionary F and to decide which algorithms we will use for sparse recovery. In practice, the dictionaries used by neural networks will likely have special properties related to the nature of the data being encoded and the operations that need to be performed. However, for the purposes of encoding a uniformly random subset of $[N]$, one natural way to choose a dictionary is to assign each codeword f_i to a random unit-norm vector. The simplest motivation for this strategy is the fact that, so long as $d = \Omega(\ln N)$, we can expect a dictionary of N randomly chosen vectors to be a “nearly orthonormal” frame in the following sense.

Proposition 1 *For $\epsilon > 0$ and $p > 0$, let $d > 2\epsilon^{-2}(2\ln N + \ln p^{-1})$, and let $\{F_1, \dots, F_N\}$ be independent draws from a uniform distribution over the unit sphere in $\mathbb{R}^N$. Then*

$$\forall i \neq j, |\langle F_i, F_j \rangle| < \epsilon$$

with probability at least $(1 - p)$.

This justifies that, for fixed k , we need only take $N = O(\ln N)$ in order for each coefficient x_i to be retrievable directly by thresholding the linear function $\langle y, f_i \rangle$ of the code y . However, we will see that the minimum values of N needed to reliably encode a random set Y varies significantly depending on the method of decoding used.

We informally refer to decoding methods that require only a single “step” of estimation and thresholding as **one-step methods**. In contrast, the field of compressive sensing provides a number of more elaborate *iterative* methods for sparse recovery. However, many of these methods are not applicable to the very large problems faced by sparse autoencoders, where performing a training run using only a one-step method is already computationally demanding. In this work, we consider a family of relatively cheap iterative algorithms which we call **matching- k methods**.

2.1 One-Step Methods

One way to motivate the construction of one-step methods is to think of the isolated subproblem of estimating a coefficient X_i from the model

$$Y = X_i f_i + \overbrace{\sum_{j \neq i} X_j f_j}^{Z_i}. \quad (1)$$

(Here we write X and Y with capital letters to emphasize they are random variables.) If we approximate the sum Z_i of other codewords as centered Gaussian noise with covariance Σ , the inference problem for X_i becomes tractable; specifically, if we define an inner product by $\langle v, w \rangle_\Sigma = \frac{1}{2} v^T \Sigma^{-1} w$, then the posterior on X_i can be expressed in terms of the linear function

$$\lambda_i(Y) = \frac{\langle f_i, Y \rangle_\Sigma}{\|f\|_\Sigma^2}.$$

In signal processing, this is called a **matched filter** for the codeword f_i . In terms of the matched filter,

$$\ln P(Y_i = y | X = x) = \frac{\rho}{2} (\lambda_i(x) - y)^2 + \ln P(Y_i = y) + C \quad (2)$$

where $\rho = \|f\|_\Sigma^2$ is called the “signal-to-noise ratio.” See Appendix A for a review.

When f_i are independent draws from some isotropic distribution, it is natural to take Σ to be some multiple of the identity, so that a maximum likelihood decoder has the form

$$\hat{X}_i(Y) = \begin{cases} 1 & : \langle f_i, Y \rangle \geq \tau \\ 0 & : \text{otherwise} \end{cases}$$

for some $\tau \in (0, 1)$. If the number k of non-zero latents is known or can be guessed in advance, one simple improvement is to take $\hat{X}$ to be k symbols that maximize the filters $\langle f_i, Y \rangle$. This is called the top- k operation.

One-step methods for the model selection problem have been studied e.g. by Bajwa et al. (2010). To extend one-step methods to the problem of sparse recovery for a vector Y whose non-zero coefficients may be arbitrary numbers, there are at least two possibilities. The first, considered by Bajwa et al. (2010), is to first identify the support of Y and then solve

, and was introduced in Makhzani & Frey (2014). Gao et al. (2024) showed that, in practice, top- k autoencoders perform better than their ReLU variants in practice.

2.2 Matching- k

Some of these are based on minimizing an objective function like

$$\|Fx - y\|_2^2 + \lambda \|x\|_1$$

for some $\lambda > 0$, which in statistics is known as lasso regression.

3 Information Theory Bounds

Suppose that the latent variable Y to be encoded is a symbol drawn from an alphabet of size N . What is the minimum dimension d such that each value of Y can be represented by a vector $X \in \mathbb{R}^d$?

When the vector X can be stored exactly, the answer to our question simply depends on the number of values it can take. If the coefficients of X are 16 bit floating point numbers, then each coefficient can store (nearly) 2^{16} distinct numbers, and overall we need only about $d_{\min} \approx \frac{1}{16} \log_2 N$ dimensions.

However, we expect that activation vectors within real neural networks are subject to various kinds of interference. For example, an activation vector X may need to represent the symbol Y while also holding information on another latent quantity, or it may be subject to explicitly induced noise like dropout. Characterizing the “capacity” of a vector in the presence of interference is a basic problem in coding theory.




Figure 3: **Solid line:** effective inverse capacity guaranteed for top- k decoding by Proposition 4, graphed as a function of $\eta = \ln k / \ln N$. **Points:**

One common model is to consider that X is subject to additive white Gaussian noise. More specifically, we let $Z \in \mathbb{R}^d$ be a vector of independent, centered Gaussians each of variance P , and consider the problem of recovering Y from $X + Z$. The following is then a typical (but remarkable) result of information theory.

Proposition 2 *Let the random variable $Y \in [N]$ be uniformly distributed and let Z be white Gaussian noise of variance P . Suppose there exist a pair of maps $F: [N] \rightarrow \mathbb{R}^d$ and $G: \mathbb{R}^d \rightarrow [N]$ so that*

$$G(F(Y) + Z) = Y$$

with probability at least $(1 - p)$, and suppose the variance of each coordinate of $F(Y)$ is bounded by 1. Then

$$d \geq 2 \frac{(1 - p) \ln N - \ln 2}{\ln(1 + P^{-1})}.$$

Asymptotically for large N , this means that

$$\frac{d}{H} \geq (1 + o(1))C(P)$$

where $H = \ln N$ is the entropy of Y and $C(P) = 2 / \ln(1 + P^{-1})$ is called the inverse **channel capacity**. For large P , $C(P) \approx 2P$. Overall, we informally conclude that we need $2P$ dimensions to store each nat of information present in Y .

Conversely, it can be shown that our bound on the ratio d/H is asymptotically tight. More concretely, we can say that for any $\epsilon > 0$ and for sufficiently large N , it is enough to take

$$\frac{d}{H} \geq (1 + \epsilon)C(P)$$

for there to exist some pair (F, G) of encoding and decoding maps that satisfy $G(F(Y) + Z) = Y$ with probability arbitrarily close to 1. For a reference on this fact, see Chapter 9 of Cover & Thomas (2006).

In the remainder of this work, we are interested in applying a similar point of view to our toy superposition coding problem. Although we do not consider the influence of noise, we will see that the minimum dimension d required for a superposition to store a k -element subset of $[N]$ can be described, in a certain asymptotic regime, as

3.1 Results for Random Dictionaries

We now return to our toy scenario. Let F be a dictionary of N codewords drawn independently and uniformly from the unit sphere in $\mathbb{R}^d$. Let $b(d, k, N)$ denote the probability that a superposition code for a k -element is *not* accurately recovered by top- k decoding.




Figure 4: The prediction of Proposition 4 agrees well with numerical experiments.

Our first result provides an asymptotic constraint on the dimension d required for $b(d, k, N)$ to be small.

Proposition 3 *Let $\epsilon > 0$ be arbitrary. Over the regime $d \leq (1 - \epsilon)2k \ln N$,*

$$\limsup_{N \rightarrow \infty} b(d, k, N) > 0.$$

Here:

- Justification

Proposition 4 *For $\eta \in (0, 1)$, let $\ln k \leq \eta \ln N$, and let*

$$C(\eta) = \frac{4}{1 - \sqrt{\eta}} - 2, \quad H = \ln \binom{N}{k}.$$

Let $\epsilon > 0$ be arbitrary. Then over the regime $d \geq (1 + \epsilon)C(\eta)H$,

$$\limsup_{N \rightarrow \infty} b(d, k, N) = 0.$$

A simple rule of thumb in practice is that around

$$d \approx 4k \ln(Nk).$$

Numerical values of the function $C(\eta)$ are given in Section 3.1.

Here: explain the new figures.

Matching pursuit far outperforms top- k decoding for the range of N and k considered earlier. When $d \geq 1.3 \log_2(eN/k)$ and $N > 2^{16}$, our experiments show that matching pursuit is very reliable. In other words, matching pursuit can reliably infer a sparse vector from just 1.3 dimensions per bit. In ?? we find that a more computationally expensive algorithm called basis pursuit can do slightly better, requiring only 0.8 dimensions per bit.

3.2 Are Random Dictionaries Optimal?

In ??, we considered the performance of threshold and top- k decodings at recovering a subset from a superposition code with a *random* dictionary F . One natural question is whether these one-step decodings can do better if the dictionary is optimized to reduce the scale of “crosstalk” between distinct codewords.




Figure 5: Empirically estimated probabilities of decoding a k -element subset of $[N]$ from a d -dimensional superposition code with random dictionary, .

Of course, when $d \geq N$, we can make the codewords f_i exactly orthogonal. For this reason, the performance of one-step decodings shown in the left-most column of ?? is much worse than is possible; we never need more than N dimensions to store a latent vector of dimension N . However, when the ratio d/N is small—say, smaller than $1/10$ —we conjecture that optimizing the dictionary gives practically no improvement over a random initialization. Unfortunately, we are not aware of a theoretical justification for this fact.

To understand our conjecture, recall the “crosstalk” terms

$$\xi_i = \sum_{j \neq i} Y_j \langle f_i, f_j \rangle.$$

For each $i \in [N]$, this is a sum of between k and $(k-1)$ numbers drawn without replacement from the sequence

$$(\langle f_i, f_j \rangle)_{j \neq i}.$$

Let’s fix the dictionary F and consider the empirical distribution defined by this sequence of $N-1$ numbers. Suppose this distribution has zero mean and variance

$$\gamma_i(F) = \frac{1}{N-1} \sum_{j \neq i} \langle f_i, f_j \rangle^2.$$

When k is moderately large but much smaller than N , we expect the crosstalk ξ_i to behave like a centered Gaussian with variance $k\gamma_i$. Specifically, we expect that the probability of its tail events with respect to the random set Y will be governed by the product $k\gamma_i$. If we assume that tail events for the different variables ξ_i are “sufficiently independent,” we conclude overall that the typical value of $\gamma_i(F)$ is the limiting factor for the reliability of one-step estimates.

A dictionary chosen to make the quantities γ_i uniformly smaller would, in particular, have smaller average squared interference

$$\gamma(F) = \frac{1}{N} \sum_{i=1}^N \gamma_i = \binom{n}{2}^{-1} \sum_{i \neq j} \langle f_i, f_j \rangle^2$$

between distinct codewords. For a random dictionary F , $\gamma(F)$ equals $1/d$ in expectation. Can we decrease this value significantly by optimization?

Using projected gradient descent, we minimized $\gamma(F)$ subject to the constraint of maintaining unit norm codewords. We tested dictionaries with between $N = 64$ and $N = 2^{16} = 65536$ codewords and with codeword dimensions between $d = 16$ and 1024 . In each case, we initialized with a random Rademacher dictionary and optimized to convergence with standard criteria. Our results are shown in Section 3.2 of ??.

As d approaches N , we find that the optimal value γ_{opt} of $\gamma(F)$ converges to 0, as expected. On the other hand, when $d \ll N$, γ_{opt} is very close to $1/d$, its expected value under a random initialization. For example, with $N = 2^{16}$ (not plotted), the optimal value of $\gamma(F)$ is indistinguishable from $1/d$ on a log-log plot.

Furthermore, we find a striking regularity. Empirically, the ratio $\gamma_{\text{opt}}/d^{-1} = d\gamma_{\text{opt}}$ between the optimal value of γ and its expected value at initialization turns out to be a function of the relative dimension d/N . Since this holds as N ranges over several orders of magnitude, it is natural to believe it may hold in general.

Claim 1 *For given (N, d) , the optimal value of $\gamma(F)$ for a dictionary $F \in \mathbb{R}^{d \times N}$ of unit norm codewords is*

$$\gamma_{\text{opt}}(N, d) = \frac{\kappa(d/N)}{d}$$

for some function κ . Furthermore, $\kappa(r)$ is close to 1 for small values of r .

If true, this means that the values $\gamma_i(F)$ governing the scale of crosstalk suffered by matched filters can’t be made significantly smaller than $1/d$ when $d \leq \epsilon N$ for small ϵ .

We’re not aware of theoretical results in this direction. Note in particular that this is not obviously related to work on sphere packing (see Cohn & Zhao (2014)) since we are interested in the *scale* of the distribution of inner products rather than in maximum values.




Figure 6: **Top:** The mean squared interference $\gamma(F)$ of a dictionary F obtained by running projected gradient descent to convergence. The dotted line shows $\gamma_{\text{init}} = 1/d$, the mean squared interference attained in expectation by a random initialization. The best interference for $N = 2^{16}$ found by gradient descent (not graphed) is nearly indistinguishable from the dotted line. **Bottom:** A plot of ratio $\gamma_{\text{opt}}/\gamma_{\text{init}}$ by which gradient descent improves γ relative to its expected value at initialization against the ratio d/N between code-word dimension and dictionary size.

References

- Waheed U. Bajwa, Robert Calderbank, and Sina Jafarpour. Why Gabor frames? Two fundamental measures of coherence and their role in model selection. *Journal of Communications and Networks*, 12(4):289–307, 2010. URL <https://ieeexplore.ieee.org/abstract/document/6388466/>. Publisher: KICS.
- Bart Bussmann, Noa Nabeshima, Adam Karvonen, and Neel Nanda. Learning Multi-Level Features with Matryoshka Sparse Autoencoders, March 2025. URL <http://arxiv.org/abs/2503.17547>. arXiv:2503.17547 [cs].
- Henry Cohn and Yufei Zhao. Sphere packing bounds via spherical codes. *Duke Mathematical Journal*, 163(10), July 2014. ISSN 0012-7094. doi: 10.1215/00127094-2738857. URL <http://arxiv.org/abs/1212.5966>. arXiv:1212.5966 [math].
- Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory 2nd Edition*. Wiley-Interscience, Hoboken, N.J, 2006. ISBN 978-0-471-24195-9.
- Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse Autoencoders Find Highly Interpretable Features in Language Models, October 2023. URL <http://arxiv.org/abs/2309.08600>. arXiv:2309.08600.
- M. Elad. *Sparse and redundant representations: from theory to applications in signal and image processing*. Springer, New York, 2010. ISBN 978-1-4419-7010-7 978-1-4419-7011-4. OCLC: ocn646114450.
- Joshua Engels, Logan Riggs, and Max Tegmark. Decomposing The Dark Matter of Sparse Autoencoders, October 2024. URL <http://arxiv.org/abs/2410.14670>. arXiv:2410.14670.
- Manaal Faruqui, Yulia Tsvetkov, Dani Yogatama, Chris Dyer, and Noah A. Smith. Sparse Overcomplete Word Vector Representations. In Chengqing Zong and Michael Strube (eds.), *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International*

- Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pp. 1491–1500, Beijing, China, July 2015. Association for Computational Linguistics. doi: 10.3115/v1/P15-1144. URL <https://aclanthology.org/P15-1144/>.
- Leo Gao, Tom Dupré la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike, and Jeffrey Wu. Scaling and evaluating sparse autoencoders, June 2024. URL <http://arxiv.org/abs/2406.04093>. arXiv:2406.04093 [cs].
- Alireza Makhzani and Brendan Frey. k-Sparse Autoencoders, March 2014. URL <http://arxiv.org/abs/1312.5663>. arXiv:1312.5663.
- Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. *Advances in Neural Information Processing Systems*, 35:17359–17372, 2022.
- Anh Nguyen, Jason Yosinski, and Jeff Clune. Multifaceted Feature Visualization: Uncovering the Different Types of Features Learned By Each Neuron in Deep Neural Networks, May 2016. URL <http://arxiv.org/abs/1602.03616>. arXiv:1602.03616 [cs].
- Chris Olah, Nick Cammarata, Ludwig Schubert, Gabriel Goh, Michael Petrov, and Shan Carter. Zoom In: An Introduction to Circuits. *Distill*, 5(3):10.23915/distill.00024.001, March 2020. ISSN 2476-0757. doi: 10.23915/distill.00024.001. URL <https://distill.pub/2020/circuits/zoom-in>.
- Senthooran Rajamanoharan, Arthur Conmy, Lewis Smith, Tom Lieberum, Vikrant Varma, Jānos Kramár, Rohin Shah, and Neel Nanda. Improving Dictionary Learning with Gated Sparse Autoencoders, April 2024. URL <http://arxiv.org/abs/2404.16014>. arXiv:2404.16014.
- David E. Rumelhart, James L. McClelland, and AU. *Parallel Distributed Processing: Explorations in the Microstructure of Cognition: Foundations*. The MIT Press, 1986. ISBN 978-0-262-29140-8. doi: 10.7551/mitpress/5236.001.0001.
- Lee Sharkey, Bilal Chughtai, Joshua Batson, Jack Lindsey, Jeff Wu, Lucius Bushnaq, Nicholas Goldowsky-Dill, Stefan Heimersheim, Alejandro Ortega, Joseph Bloom, Stella Biderman, Adria Garriga-Alonso, Arthur Conmy, Neel Nanda, Jessica Rumbelow, Martin Wattenberg, Nandi Schoots, Joseph Miller, Eric J. Michaud, Stephen Casper, Max Tegmark, William Saunders, David Bau, Eric Todd, Atticus Geiger, Mor Geva, Jesse Hoogland, Daniel Mufet, and Tom McGrath. Open Problems in Mechanistic Interpretability, January 2025. URL <http://arxiv.org/abs/2501.16496>. arXiv:2501.16496 [cs].
- Adly Templeton, Tom Conerly, Jonathan Marcus, Jack Lindsey, Trenton Bricken, Brian Chen, Adam Pearce, Craig Citro, Emmanuel Ameisen, Andy Jones, Hoagy Cunningham, Nicholas L Turner, Callum McDougall, Monte MacDiarmid, C. Daniel Freeman, Theodore R. Sumers, Edward Rees, Joshua Batson, Adam Jermy, Shan Carter, Chris Olah, and Tom Henighan. Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet. *Transformer Circuits Thread*, May 2024. URL <https://transformer-circuits.pub/2024/scaling-monosemanticity/index.html>.
- Simon Thorpe. Local vs. Distributed Coding. *Intellectica*, 8(2):3–40, 1989. doi: 10.3406/intel.1989.873. URL https://www.persee.fr/doc/intel_0769-4113_1989_num_8_2_873. Publisher: Persée - Portail des revues scientifiques en SHS.
- Zeyu Yun, Yubei Chen, Bruno Olshausen, and Yann LeCun. Transformer visualization via dictionary learning: contextualized embedding as a linear superposition of transformer factors. In Eneko Agirre, Marianna Apidianaki, and Ivan Vulić (eds.), *Proceedings of Deep Learning Inside Out (DeeLIO): The 2nd Workshop on Knowledge Extraction and Integration for Deep Learning Architectures*, pp. 1–10, Online, June 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.deelio-1.1. URL <https://aclanthology.org/2021.deelio-1.1/>.
- Changwan Zhang and Yue Wang. A sample survey study of poly-semantic neurons in deep CNNs. In *International Conference on Computer Graphics, Artificial Intelligence, and Data Processing (ICCAID 2022)*, volume 12604, pp. 849–855. SPIE, May 2023. doi: 10.1117/12.2674650.

A Matched Filters

Consider the problem of inferring a scalar S from the sum

$$X = Sf + Z$$

where $f \in \mathbb{R}^n$ and Z is a Gaussian variable independent from S . Suppose for simplicity that Z has non-singular covariance Σ , so that $-\ln p(z) = 1/2 \|z\|_\Sigma^2$ where

$$\|z\|_\Sigma^2 = z^T \Sigma^{-1} z.$$

Then a routine calculation shows that

$$\begin{aligned} -\ln p(S = s | X = x) \\ = C(x) - \ln p(s) + \frac{1}{2} \left(s - \frac{\langle f, x \rangle_\Sigma}{\|f\|_\Sigma^2} \right)^2 \|f\|_\Sigma^2 \end{aligned} \quad (3)$$

where $C(x)$ is a constant depending only on x and $\langle -, - \rangle_\Sigma$ is the inner product associated with the norm $\|-\|_\Sigma$. In particular, the distribution of S conditional on X is only a function of the inner product $\langle f, X \rangle_\Sigma$. The **matched filter** for S is the linear function

$$\lambda(X) = \frac{\langle f, X \rangle_\Sigma}{\|f\|_\Sigma^2},$$

and can be understood as providing the maximum likelihood estimate for S .

The quality of our matched filter is measured by its signal-to-noise ratio (SNR)

$$\rho = \frac{(\lambda(f))^2}{\text{Var}_Z \lambda(Z)} = \|f\|_\Sigma^2.$$

Up to a scalar, λ can be characterized as the linear function that maximizes this quantity. Given an improper uniform prior on S , Equation (3) shows the posterior on S conditional on X is a Gaussian with mean $\lambda(X)$ and precision ρ .

B Estimates for the Binomial Coefficient

To estimate $\ln \binom{N}{k}$, it is helpful to first remember the elementary inequalities

$$\left(\frac{N}{k} \right)^k \leq \binom{N}{k} \leq \left(\frac{eN}{k} \right)^k.$$

Taking logarithms gives

$$k \ln \left(\frac{N}{k} \right) \leq \ln \binom{N}{k} \leq k \ln \left(\frac{eN}{k} \right),$$

and so $\ln \binom{N}{k} = k \ln \left(\frac{N}{k} \right) + O(k)$. In the sublinear regime $\ln k \sim \eta \ln N$, this proves that

$$\ln \binom{N}{k} \sim (1 - \eta) k \ln N.$$

This is the approximation we need for our asymptotic results. However, to estimate for $\ln \binom{N}{k}$ when $k^2 \ll N$, we can use the finer asymptotic

$$\ln \binom{N}{k} = k \ln \left(\frac{eN}{k} \right) + O(k^2/N).$$

If we put $\eta = \ln k / \ln N$ and assume $\eta < 1/2$, this means that

$$\ln \binom{N}{k} = (1 - \eta) k \ln N + k + o(1)$$

for large N .

One heuristic justification for this finer estimate is to consider a random subset $Y \subseteq [N]$ where each element is included independently with probability $s = k/N$. Then, where

$$h(s) = -s \ln s - (1-s) \ln(1-s) = -s \ln s + s + O(s^2)$$

is the binary entropy function, the entropy of Y is

$$\begin{aligned} H(Y) &= h(s)N = -sN \ln s + sN + O(s^2 N) \\ &= k \ln \left(\frac{eN}{k} \right) + O \left(\frac{k^2}{N} \right). \end{aligned}$$

We expect that for large N the number of elements in Y concentrates sharply around k , and so $H(Y)$ should be close to $\ln \binom{N}{k}$.

Indeed this is true. One simple approach is to use Stirling's approximation

$$\ln n! = n \ln n - n + \frac{1}{2} \ln(2\pi n) + O(n^{-1})$$

to derive

$$\begin{aligned} \ln \binom{N}{k} &= \ln N! - \ln k! - \ln(N-k)! \\ &= k \ln \frac{N}{k} + (N-k) \ln \left(1 + \frac{k}{N-k} \right) + O \left(\frac{1}{N} \right). \end{aligned}$$

For $k \ll N$, we have

$$\ln \left(1 + \frac{k}{N-k} \right) = \frac{k}{N} + O \left(\frac{k^2}{N^2} \right).$$

Overall this proves that

$$\ln \binom{N}{k} = k \ln \left(\frac{eN}{k} \right) + O \left(\frac{k^2}{N} \right).$$